

Cs231n Lecture 4


Introduction to Neural Networks

CONTENTS

1. Backpropagation
2. Neural Networks
3. Artificial Neural Network

1. Backpropagation

1. Backpropagation: a simple example



1. Backpropagation: a simple example

Chain Rule : 합성함수의 미분에서 사용하는 공식


$$[f(g(x))]' = f'(g(x)) \times g'(x)$$

$$\frac{df}{dx} = \frac{df}{dg} \times \frac{dg}{dx}$$


$$f(g(h(x)))' \rightarrow \frac{df}{dx} = \frac{df}{dg} \times \frac{dg}{dh} \times \frac{dh}{dx}$$

1. Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$




1. Backpropagation: a simple example



$$f(x, y, z) = (x + y)z$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

1. Backpropagation: a simple example



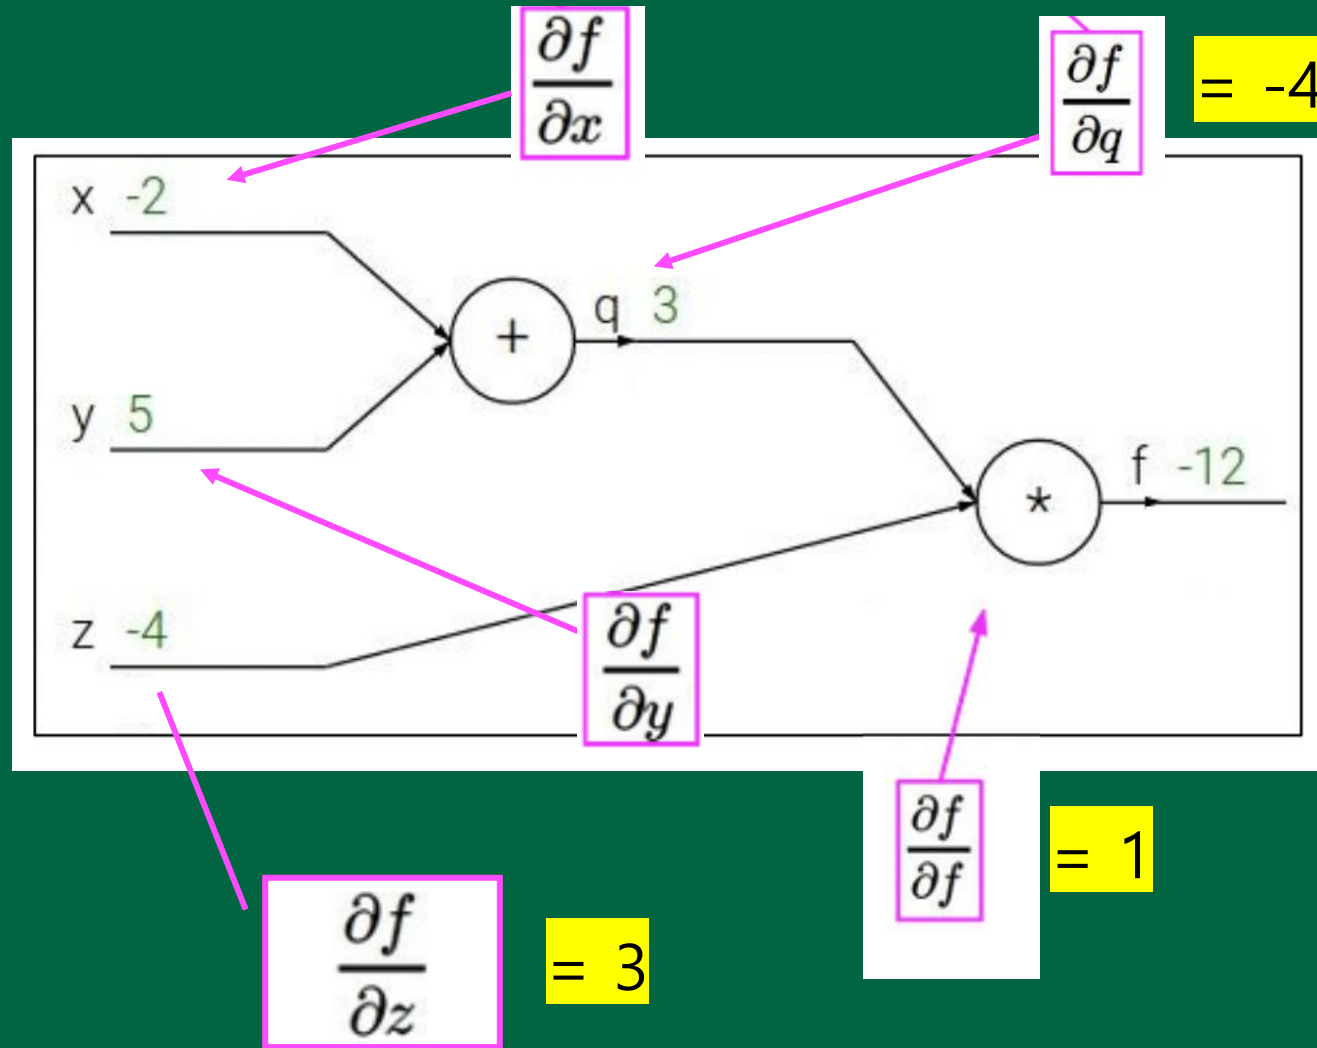
Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

$$f(x, y, z) = (x + y)z$$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

1. Backpropagation: a simple example

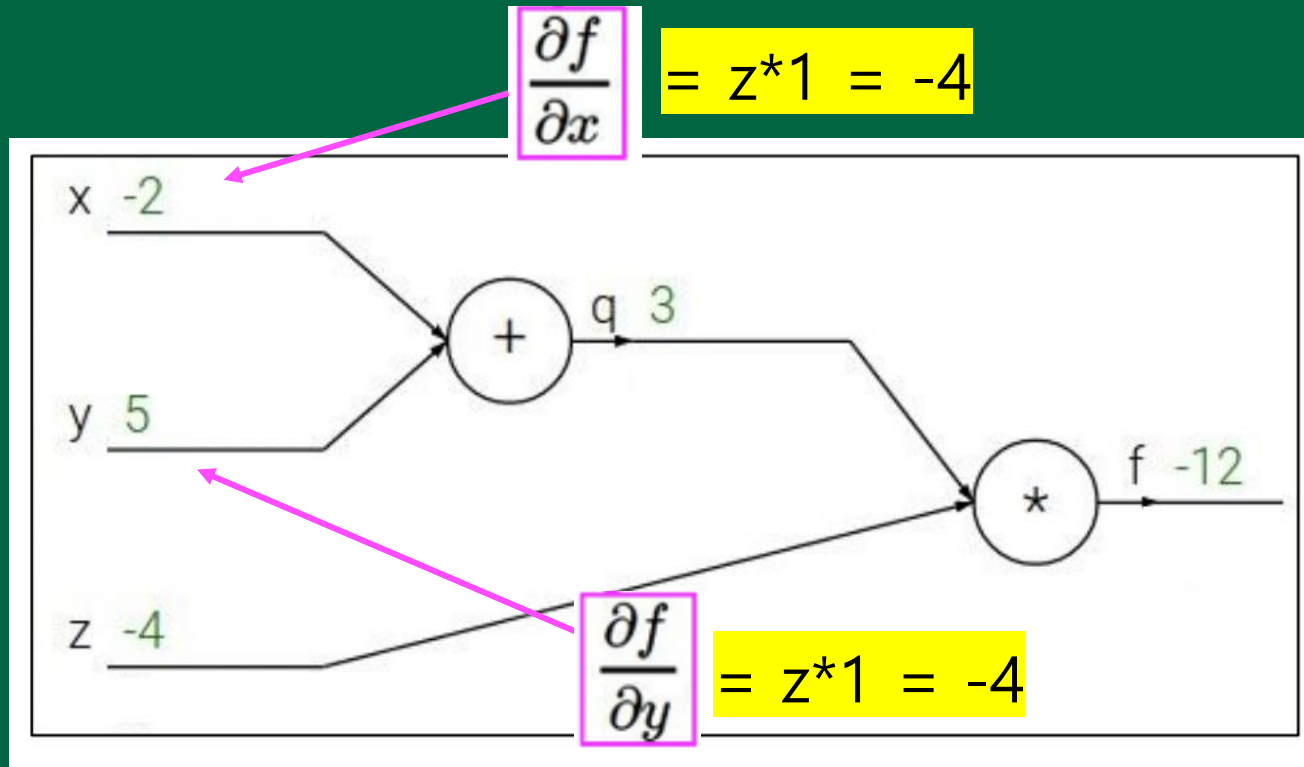


$$f(x, y, z) = (x + y)z$$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

1. Backpropagation: a simple example




Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

1. Backpropagation: a simple example




1. Backpropagation: a simple example

Remember:


$$\text{Gradient} = \text{Local Gradient} * \text{Global Gradient}$$

1. Backpropagation: another example

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$



1. Backpropagation: another example




$$f(x) = \frac{1}{x}$$

\rightarrow

$$\frac{df}{dx} = -1/x^2$$

$$-\frac{1}{1.37^2} * (1.00) = -0.53$$

1. Backpropagation: another example



Green color text : local gradient
Red color text: upstream gradient

1. Backpropagation: another example


$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$


sigmoid function

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{(1 + e^{-x})^2} = \left(\frac{1 + e^{-x} - 1}{1 + e^{-x}} \right) \times \left(\frac{1}{1 + e^{-x}} \right) = (1 - \sigma(x))\sigma(x)$$

1. Backpropagation: another example



1. Backpropagation: Patterns in backward flow




Add gate: gradient distributor

Max gate: gradient router


Mul gate: gradient switcher

1. Backpropagation: Gradients add at branches



1. Backpropagation: Gradients for vectorized code

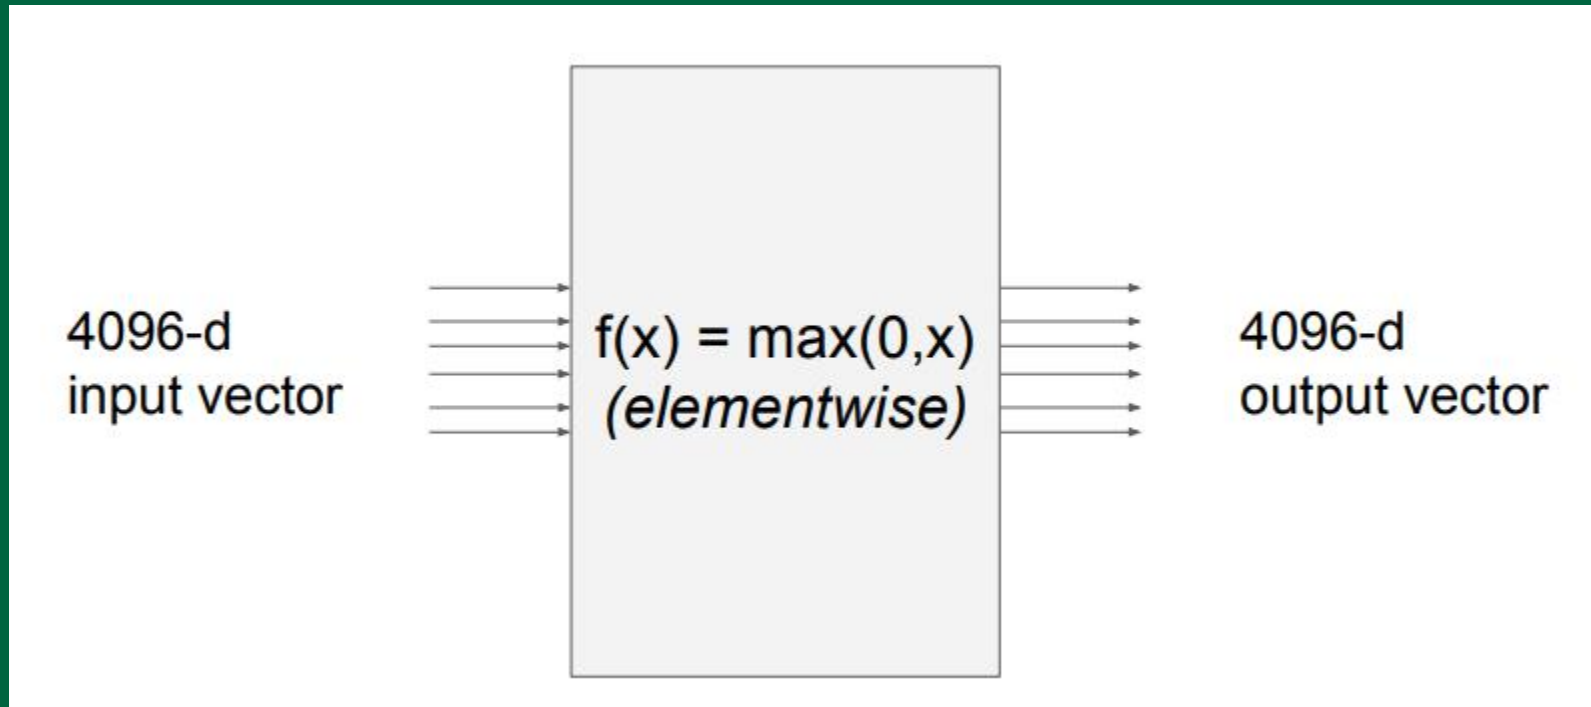
Gradients for vectorized code



1. Backpropagation: Gradients for vectorized code


$$J = \frac{\partial(x, y)}{\partial(u, v)} = \begin{vmatrix} \frac{\partial x}{\partial u} & \frac{\partial x}{\partial v} \\ \frac{\partial y}{\partial u} & \frac{\partial y}{\partial v} \end{vmatrix} = \frac{\partial x}{\partial u} \frac{\partial y}{\partial v} - \frac{\partial x}{\partial v} \frac{\partial y}{\partial u}.$$

1. Backpropagation: Vectorized operations



Jacobian Matrix size : 4096 x 4096!

1. Backpropagation: Vectorized operations



Jacobian Matrix size(100 minibatch) : 409600 x 409600!

1. Backpropagation: Vectorized operations


$$\begin{pmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 0 & 5 \end{pmatrix}$$

Diagonal Matrix

1. Backpropagation: A Vectorized example


A vectorized example: $f(x, W) = \|W \cdot x\|^2 = \sum_{i=1}^n (W \cdot x)_i^2$

$\downarrow \quad \downarrow$
 $\in \mathbb{R}^n \quad \in \mathbb{R}^{n \times n}$



1. Backpropagation: Vectorized operations

A vectorized example: $f(x, W) = ||W \cdot x||^2 = \sum_{i=1}^n (W \cdot x)_i^2$




$$q = W \cdot x = \begin{pmatrix} W_{1,1}x_1 + \cdots + W_{1,n}x_n \\ \vdots \\ W_{n,1}x_1 + \cdots + W_{n,n}x_n \end{pmatrix}$$

$$f(q) = ||q||^2 = q_1^2 + \cdots + q_n^2$$

$$\frac{\partial f}{\partial q_i} = 2q_i$$

$$\begin{aligned} \frac{\partial f}{\partial W_{i,j}} &= \sum_k \frac{\partial f}{\partial q_k} \frac{\partial q_k}{\partial W_{i,j}} \\ &= \sum_k (2q_k) (\mathbf{1}_{k=i} x_j) \\ &= 2q_i x_j \end{aligned}$$


1. Backpropagation: code example



Graph (or Net) object *(rough psuedo code)*

```
class ComputationalGraph(object):  
    #...  
    def forward(inputs):  
        # 1. [pass inputs to input gates...]  
        # 2. forward the computational graph:  
        for gate in self.graph.nodes_topologically_sorted():  
            gate.forward()  
        return loss # the final gate in the graph outputs the loss  
    def backward():  
        for gate in reversed(self.graph.nodes_topologically_sorted()):  
            gate.backward() # little piece of backprop (chain rule applied)  
        return inputs_gradients
```

1. Backpropagation: code example



(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        return z  
    def backward(dz):  
        # dx = ... #todo  
        # dy = ... #todo  
        return [dx, dy]
```


$$\frac{\partial L}{\partial z}$$

An arrow points from this box to the `dz` parameter in the `backward` method of the `MultiplyGate` class.

$$\frac{\partial L}{\partial x}$$

An arrow points from this box to the `dx` element in the `[dx, dy]` return list of the `backward` method of the `MultiplyGate` class.

1. Backpropagation: code example




(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        self.x = x # must keep these around!  
        self.y = y  
        return z  
    def backward(dz):  
        dx = self.y * dz # [dz/dx * dL/dz]  
        dy = self.x * dz # [dz/dy * dL/dz]  
        return [dx, dy]
```

2. Neural Networks

2. Neural Networks: without the brain stuff

(Before) Linear score function: $f = Wx$




$$X = 3072 \times 1$$

$$W = 10 \times 3072$$

$$S(\text{score}) = 10 \times 1$$

2. Neural Networks: without the brain stuff

(**Now**) 2-layer Neural Network $f = W_2 \max(0, W_1 x)$



$$W1 = 100 \times 3072$$

$$W2 = 10 \times 100$$


$$f = W_3 \max(0, W_2 \max(0, W_1 x))$$

2. Neural Networks: code example


```
1  import numpy as np
2  from numpy.random import randn
3
4  N, D_in, H, D_out = 64, 1000, 100, 10
5  x, y = randn(N, D_in), randn(N, D_out)
6  w1, w2 = randn(D_in, H), randn(H, D_out)
7
8  for t in range(2000):
9      h = 1 / (1 + np.exp(-x.dot(w1)))
10     y_pred = h.dot(w2)
11     loss = np.square(y_pred - y).sum()
12     print(t, loss)
13
14     grad_y_pred = 2.0 * (y_pred - y)
15     grad_w2 = h.T.dot(grad_y_pred)
16     grad_h = grad_y_pred.dot(w2.T)
17     grad_w1 = x.T.dot(grad_h * h * (1 - h))
18
19     w1 -= 1e-4 * grad_w1
20     w2 -= 1e-4 * grad_w2
```

3. Artificial Neural Network


3. Neural Networks



3. Neural Networks



3. Neural Networks: Activation functions



sigmoid activation function


$$\frac{1}{1 + e^{-x}}$$

```
class Neuron:
    # ...
    def neuron_tick(inputs):
        """ assume inputs and weights are 1-D numpy arrays and bias is a number """
        cell_body_sum = np.sum(inputs * self.weights) + self.bias
        firing_rate = 1.0 / (1.0 + math.exp(-cell_body_sum)) # sigmoid activation func
        return firing_rate
```

3. Neural Networks: Activation functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



3. Neural Networks: Architectures



“2-layer Neural Net”, or
“1-hidden-layer Neural Net”



“3-layer Neural Net”, or
“2-hidden-layer Neural Net”

“Fully-connected” layers

3. Neural Networks: code example

```
class Neuron:
    # ...
    def neuron_tick(inputs):
        """ assume inputs and weights are 1-D numpy arrays and bias is a number """
        cell_body_sum = np.sum(inputs * self.weights) + self.bias
        firing_rate = 1.0 / (1.0 + math.exp(-cell_body_sum)) # sigmoid activation function
        return firing_rate
```


THANK YOU